

DavisSML Tensorflow Lab 1

Some content is taken from

- www.tensorflow.org (<http://www.tensorflow.org>)
- <http://rail.eecs.berkeley.edu/deeprlcourse/> (<http://rail.eecs.berkeley.edu/deeprlcourse/>)

To install tensorflow use pip: `pip install tensorflow`

There are 2 exercises below. Do not use Keras or another high level API.

Tensors

```
In [1]: 1 import tensorflow as tf
        2 import numpy as np
        3 import matplotlib.pyplot as plt
```

- Tensorflow is based around tensor objects
- Tensors can be symbolic or actual data
- Eager execution actually creates the tensors, not just symbolically. Without it the Tensors would only be symbolic until a session is started.

```
In [2]: 1 tf.enable_eager_execution()
```

Several types of special tensors are

- `tf.Variable` : typically parameters in your model
- `tf.constant` : fixed values (gradients will not be calculated wrt them)
- `tf.placeholder` : a tensor waiting to be assigned a value; not compatible with eager execution
- `tf.SparseTensor` : a sparse tensor

```
In [3]: 1 # initialize these constants
        2 a = tf.constant(1.0)
        3 b = tf.constant(2.0)
        4
        5 # mathematical operations are defined on tensors
        6 c = a + b
        7
        8 print(c)
        tf.Tensor(3.0, shape=(), dtype=float32)
```

Above are scalars which are order 0 tensors, other types are below:

```
In [4]: 1 mammal = tf.Variable("Elephant", tf.string) # order 0
        2 ignition = tf.Variable(451, tf.int16)
        3 floating = tf.Variable(3.14159265359, tf.float64)
        4 its_complicated = tf.Variable(12.3 - 4.85j, tf.complex64)
```

```
In [5]: 1 mammal.get_shape()
```

```
Out[5]: TensorShape ([])
```

Here are some order 1 tensors of various types:

```
In [6]: 1 mystr = tf.Variable(["Hello", "World"], tf.string) #order 1
2 cool_numbers = tf.Variable([3.14159, 2.71828], tf.float32)
3 first_primes = tf.Variable([2, 3, 5, 7, 11], tf.int32)
4 its_very_complicated = tf.Variable([12.3 - 4.85j, 7.5 - 6.23j], tf.complex64)
```

The shape of the tensor is like that of numpy arrays, they will tell you the order and dimensions.

```
In [7]: 1 its_complicated.get_shape(), its_very_complicated.get_shape(), first_primes.get_shape
Out[7]: (TensorShape([]), TensorShape([Dimension(2)]), TensorShape([Dimension(5)]))
```

```
In [8]: 1 beta_init = np.random.uniform(0,1, (32,32,32))
2 betas = [tf.convert_to_tensor(beta_init[0,0,0], dtype=tf.float32),
3          tf.convert_to_tensor(beta_init[:,0,0], dtype=tf.float32),
4          tf.convert_to_tensor(beta_init[:, :,0], dtype=tf.float32), #order 2
5          tf.convert_to_tensor(beta_init, dtype=tf.float32)]
6
7 [beta.get_shape() for beta in betas]
Out[8]: [TensorShape([]),
TensorShape([Dimension(32)]),
TensorShape([Dimension(32), Dimension(32)]),
TensorShape([Dimension(32), Dimension(32), Dimension(32)])]
```

Tensor operations input and output tensors, and will cast to a tensor if needed.

```
In [9]: 1 print(tf.square(5))
2 print(tf.reduce_sum([1, 2, 3]))
tf.Tensor(25, shape=(), dtype=int32)
tf.Tensor(6, shape=(), dtype=int32)
```

Even when in eager mode Tensors are different than arrays because

- they are immutable : the data cannot be modified without reassigning the variable
- they can live on GPUs as comfortably as CPUs

Why are they immutable?

You can assign slices to other variables:

```
In [10]: 1 X = tf.Variable([[1,2,3],[3,4,5]])
2 Y = X[:,2,:2]
```

```
In [11]: 1 Y
```

```
Out[11]: <tf.Tensor: id=96, shape=(2, 2), dtype=int32, numpy=
array([[1, 2],
       [3, 4]])>
```

```
In [12]: 1 X.assign([[2,2,3],[3,4,5]])
```

```
Out[12]: <tf.Variable 'UnreadVariable' shape=(2, 3) dtype=int32, numpy=
array([[2, 2, 3],
       [3, 4, 5]])>
```

```
In [13]: 1 X, Y
```

```
Out[13]: (<tf.Variable 'Variable:0' shape=(2, 3) dtype=int32, numpy=
array([[2, 2, 3],
       [3, 4, 5]])>, <tf.Tensor: id=96, shape=(2, 2), dtype=int32, numpy=
array([[1, 2],
       [3, 4]])>)
```

```
In [14]: 1 beta = tf.convert_to_tensor(beta_init, dtype=tf.float32)
          2 beta.device
Out[14]: '/job:localhost/replica:0/task:0/device:CPU:0'
```

Can use GPU if it is available, because many operations are then parallel, it can be much faster:

```
In [16]: 1 import time
          2
          3 def time_matmul(x):
          4     start = time.time()
          5     for loop in range(10):
          6         tf.matmul(x, x)
          7
          8     result = time.time() - start
          9
         10     print("10 loops: {:.2f}ms".format(1000*result))
         11
         12
         13 # Force execution on CPU
         14 print("On CPU:")
         15 with tf.device("CPU:0"):
         16     x = tf.random_uniform([2000, 2000])
         17     assert x.device.endswith("CPU:0")
         18     time_matmul(x)
         19
         20 # Force execution on GPU #0 if available
         21 if tf.test.is_gpu_available():
         22     print("On GPU:")
         23     with tf.device("GPU:0"): # Or GPU:1 for the 2nd GPU, GPU:2 for the 3rd etc.
         24         x = tf.random_uniform([2000, 2000])
         25         assert x.device.endswith("GPU:0")
         26         time_matmul(x)
         27 else:
         28     print('GPU unavailable')
On CPU:
10 loops: 844.74ms
On GPU:
10 loops: 1.00ms
```

Automatic differentiation

Typically automatic differentiation works behind the scenes in tensorflow, but with `gradienttape` you can actually interact with the gradients. When you start the `gradienttape`, it will start tracking the operations performed and thus can compute gradients.

In this case the resulting gradient is also a tensor:

```
In [17]: 1  ## simple gradient
          2
          3  print("Running simple gradient ex.")
          4
          5  x = tf.ones((2, 2))
          6
          7  with tf.GradientTape() as t:
          8      t.watch(x)
          9      y = tf.reduce_sum(x)
         10      z = tf.multiply(y, y)
         11
         12  dz_dx = t.gradient(z, x)
         13
         14  print(dz_dx)
Running simple gradient ex.
tf.Tensor(
[[8. 8.]
 [8. 8.]], shape=(2, 2), dtype=float32)
```

Exercise 1. (5 pts) Below is incomplete code for logistic regression gradient computation. Fill in the necessary components.

```

In [32]: 1  ## logistic regression gradient
          2
          3  print("Running log.reg. ex.")
          4
          5  n, p = 10,5
          6  X_val = np.random.uniform(0,1,(n,p),)
          7  beta_val = np.random.uniform(0,1,(p,1))
          8  p_val = 1. / (1. + np.exp(-X_val @ beta_val))
          9  y_val = 2*(np.random.uniform(0,1,(n)) < p_val) - 1
         10
         11  ### Create X,y,beta tensors of the same shape from the above arrays
         12  X = tf.convert_to_tensor(X_val, dtype=tf.float32)
         13  y = tf.convert_to_tensor(y_val, dtype=tf.float32)
         14  beta = tf.convert_to_tensor(beta_val, dtype=tf.float32)
         15
         16
         17
         18  with tf.GradientTape(persistent=True) as t:
         19      t.watch(beta)
         20      z = tf.tensordot(y, tf.matmul(X,beta),axes=1)
         21      w = tf.reduce_sum(-tf.log(1 + tf.exp(-z)))
         22      z2 = tf.matmul(y, tf.tensordot(X,beta,axes=1))
         23      w2 = tf.reduce_sum(-tf.log(1 + tf.exp(-z2)))
         24      ### create z an order 1, with shape [10], tensor of the (y X \beta) vector.
         25      ### Matrix multiply is tf.multiply, tf.tensordot is dot product
         26
         27      ### create a tensor for the sum of - log likelihood (log(1 + exp(-z)))
         28
         29  # calculate the gradient of the loss and print it
         30  grad1 = t.gradient(w, beta)
         31  grad2 = t.gradient(w2, beta)
         32  print(grad1)
         33  print(grad2)
Running log.reg. ex.
tf.Tensor(
[[0.14849472]
 [0.30762616]
 [0.41728202]
 [0.90182316]
 [0.58966696]], shape=(5, 1), dtype=float32)
tf.Tensor(
[[0.14849472]
 [0.30762616]
 [0.41728202]
 [0.90182316]
 [0.58966696]], shape=(5, 1), dtype=float32)

```

I tried 2 methods to calculate gradients. Both gave the same answer

Sessions

We will leave eager mode, to see how sessions work.

```

In [ ]: 1  import os
          2  os._exit(00)

In [1]: 1  import tensorflow as tf
          2  import numpy as np
          3  import matplotlib.pyplot as plt

```

```
In [2]: 1 def tf_reset_graph():
2         try:
3             sess.close()
4         except:
5             pass
6         new_graph = tf.Graph()
7         return new_graph, tf.Session(graph=new_graph)
8
9 def tf_reset_default():
10        try:
11            sess.close()
12        except:
13            pass
14        tf.reset_default_graph()
15        return tf.Session()
```

```
In [3]: 1 tf_graph, sess = tf_reset_graph()
```

A computation graph contains

- tensors defined in with that graph as default
- operations that relate the tensors

The session is associated with a graph. Typically, you will be working with a single graph, so you do not have to explicitly associate the operations/tensors to a graph, and can rely on the default.

```
In [4]: 1 # define your inputs
2 with tf_graph.as_default():
3     a = tf.constant(1.0)
4     b = tf.constant(2.0)
5
6     # do some operations
7     c = a + b
8
9     # get the result
10    sess.run(c)
```

```
Out[4]: 3.0
```

Notice that the session above is still open and needs to be closed. If you look at `tf_reset`, this is what happens in the `try:` statement.

With a session it makes sense to have placeholder tensors that can be fed data. You should think of this as taking the place of a constant, where the constant value is determined at run time.

```
In [5]: 1 tf_graph, sess = tf_reset_graph()
2
3 with tf_graph.as_default():
4     # define your inputs
5     a = tf.placeholder(dtype=tf.float32, shape=[1], name='a_placeholder')
6     b = tf.placeholder(dtype=tf.float32, shape=[1], name='b_placeholder')
7
8     # do some operations
9     c = a + b
10
11 # get the result
12 c0_run = sess.run(c, feed_dict={a: [1.0], b: [2.0]})
13 c1_run = sess.run(c, feed_dict={a: [2.0], b: [4.0]})
14
15 print('c0 = {}'.format(c0_run))
16 print('c1 = {}'.format(c1_run))
c0 = [3.]
c1 = [6.]
```

You can allow the placeholder tensor to have variable dimensions by making the dimension input to be None.

```
In [6]: 1 sess = tf_reset_default()
2
3 # inputs
4 a = tf.placeholder(dtype=tf.float32, shape=[None], name='a_placeholder')
5 b = tf.placeholder(dtype=tf.float32, shape=[None], name='b_placeholder')
6
7 # do some operations
8 c = a + b
9
10 # get outputs
11 c0_run = sess.run(c, feed_dict={a: [1.0], b: [2.0]})
12 c1_run = sess.run(c, feed_dict={a: [1.0, 2.0], b: [2.0, 4.0]})
13
14 print(a)
15 print('a shape: {}'.format(a.get_shape()))
16 print(b)
17 print('b shape: {}'.format(b.get_shape()))
18 print('c0 = {}'.format(c0_run))
19 print('c1 = {}'.format(c1_run))
Tensor("a_placeholder:0", shape=(?,), dtype=float32)
a shape: (?,)
Tensor("b_placeholder:0", shape=(?,), dtype=float32)
b shape: (?,)
c0 = [3.]
c1 = [3. 6.]
```

Layers are operations that add variables to the default graph. They are used to make the code more concise and development easier. The following Dense layer will add a scalar coefficient and an intercept.

Optimizers are operations that build out the compute graph for a stochastic gradient descent algorithm. It will create the necessary tensors to make gradient steps, and can be standard SGD, RMSprop, Adam, etc.

In [7]:

```

1 sess = tf.reset_default()
2
3 ## Silly data
4 x = tf.constant([[1], [2], [3], [4]], dtype=tf.float32)
5 y_true = tf.constant([[0], [-1], [-2], [-3]], dtype=tf.float32)
6
7 """
8 This layer implements the operation: outputs = activation(inputs * coefficient + intercept)
9 Where activation is the activation function passed as the activation argument (if not None)
10 kernel (coefficient) is a weights matrix created by the layer,
11 and bias (intercept) is a bias vector created by the layer (only if use_bias is True).
12 """
13 The layer is an operation with variables (to be seen)
14 """
15 linear_model = tf.layers.Dense(units=1)
16
17 """
18 Apply the linear_model operation to tensor x to produce tensor y_pred
19 """
20
21 y_pred = linear_model(x)
22
23 """
24 Losses are output scalar tensors and can accept weights, etc.
25 """
26
27 loss = tf.losses.mean_squared_error(labels=y_true, predictions=y_pred)
28
29 """
30 Can either use a scalar learning rate or tensor (I want to change mine so I use a tensor)
31 """
32 learning_rate = tf.placeholder(tf.float32, shape=[])
33
34 """
35 train module contains optimizers such as AdaGrad and Adam
36 - contains methods: apply_gradients, compute_gradients, variables, minimize
37 - minimize returns an operation that minimizes the input (loss) tensor
38 """
39
40 optimizer = tf.train.GradientDescentOptimizer(learning_rate)
41 train = optimizer.minimize(loss)
42
43 """
44 You can initialize variables by hand or just use random starts like below
45 """
46 init = tf.global_variables_initializer()
47
48 """
49 TensorFlow uses the tf.Session class to represent a connection between the client program
50 and the server
51 - the session should be closed to free up the device, or even better, use the with clause
52 - the run method will execute the subgraph in order to evaluate the tensor or run the operation
53 """
54
55
56 sess.run(init)
57
58 for i in range(100):
59     lr = .5*(i+1)**-1.
60     _, loss_value = sess.run((train, loss), feed_dict={learning_rate : lr})
61     print(loss_value)

```


26.320143

1400.1321


```
In [15]: 1 print(sess.run(y_pred))
          2 print(sess.run(linear_model.weights))
          [-0.27734506]
          [-1.1343927 ]
          [-1.9914403 ]
          [-2.8484879 ]]
          [array([[ -0.8570476]], dtype=float32), array([0.57970256], dtype=float32)]
```

The following example from UC Berkeley CS294, trains a fully connected neural net (with 3 hidden layers) for a dataset (below) with scalar input.

```
In [16]: 1 # generate the data
          2 inputs = np.linspace(-2*np.pi, 2*np.pi, 10000)[: , None]
          3 outputs = np.sin(inputs) + 0.05 * np.random.normal(size=[len(inputs),1])
          4
          5 plt.scatter(inputs[:, 0], outputs[:, 0], s=0.1, color='k', marker='o')
Out[16]: <matplotlib.collections.PathCollection at 0x14cfff90ba8>
```



In [21]:

```

1  sess = tf.reset_default()
2
3  def create_model():
4      # create inputs
5      input_ph = tf.placeholder(dtype=tf.float32, shape=[None, 1])
6      output_ph = tf.placeholder(dtype=tf.float32, shape=[None, 1])
7
8      # create variables
9      W0 = tf.get_variable(name='W0', shape=[1,3], initializer=tf.contrib.layers.xavier_in
10     #W1 = tf.get_variable(name='W1', shape=[20, 20], initializer=tf.contrib.layers.xavie
11     #W2 = tf.get_variable(name='W2', shape=[20, 1], initializer=tf.contrib.layers.xavier
12
13     b0 = tf.get_variable(name='b0', shape=[3], initializer=tf.constant_initializer
14     #b1 = tf.get_variable(name='b1', shape=[20], initializer=tf.constant_initializer(0.)
15     #b2 = tf.get_variable(name='b2', shape=[1], initializer=tf.constant_initializer(0.))
16
17     weights = [W0]#[W0, W1, W2]
18     biases = [b0]#[b0, b1, b2]
19     activations = [tf.nn.relu]#[tf.nn.sigmoid, tf.nn.relu, tf.nn.sigmoid]
20
21     # create computation graph
22     layer = input_ph
23     for W, b, activation in zip(weights, biases, activations):
24         layer = tf.matmul(layer, W) + b
25         if activation is not None:
26             layer = activation(layer)
27     output_pred = layer
28
29     return input_ph, output_ph, output_pred
30
31 input_ph, output_ph, output_pred = create_model()
32
33 # create loss
34 mse = tf.reduce_mean(0.5 * tf.square(output_pred - output_ph))
35
36 # create optimizer
37 opt = tf.train.AdamOptimizer().minimize(mse)
38
39 # initialize variables
40 sess.run(tf.global_variables_initializer())
41 # create saver to save model variables
42 saver = tf.train.Saver()
43
44 # run training
45 batch_size = 32
46 for training_step in range(10000):
47     # get a random subset of the training data
48     indices = np.random.randint(low=0, high=len(inputs), size=batch_size)
49     input_batch = inputs[indices]
50     output_batch = outputs[indices]
51
52     # run the optimizer and get the mse
53     _, mse_run = sess.run([opt, mse], feed_dict={input_ph: input_batch, output_ph
54
55     # print the mse every so often
56     if training_step % 1000 == 0:
57         print('{0:04d} mse: {1:.3f}'.format(training_step, mse_run))
58     saver.save(sess, '/tmp/model.ckpt')
0000 mse: 2.115
1000 mse: 0.230
2000 mse: 0.238
3000 mse: 0.231
4000 mse: 0.214
5000 mse: 0.140


```

```

In [22]: 1 sess = tf_reset_default()
          2
          3 # create the model
          4 input_ph, output_ph, output_pred = create_model()
          5
          6 # restore the saved model
          7 saver = tf.train.Saver()
          8 saver.restore(sess, "/tmp/model.ckpt")
          9
         10 output_pred_run = sess.run(output_pred, feed_dict={input_ph: inputs})
         11
         12 plt.scatter(inputs[:, 0], outputs[:, 0], c='k', marker='o', s=0.1)
         13 plt.scatter(inputs[:, 0], output_pred_run[:, 0], c='r', marker='o', s=0.1)
INFO:tensorflow:Restoring parameters from /tmp/model.ckpt

```

Out[22]: <matplotlib.collections.PathCollection at 0x14d0076c0f0>



```

In [32]: 1 n = 10000
          2 X = np.random.uniform(0,1,(n,2))
          3 p = 1 / (1 + np.exp(-5*(X.sum(axis=1) - 2.* (X[:,0] * X[:,1] > .25) - .5)))
          4 y = np.random.uniform(0,1,n) < p

```

```

In [33]: 1 import plotnine as p9
          2 import pandas as pd

```

```

In [34]: 1 XOR_data = pd.DataFrame(X,columns=['X0','X1'])
          2 XOR_data['p'] = p


```

```

In [59]: 1 XOR_data.head()
          2 inputs = np.array(XOR_data.iloc[:,0:2]).astype('float32')
          3 outputs = np.array(XOR_data.iloc[:,2]).reshape((inputs.shape[0],1))
          4 print(inputdata.shape,outputs.shape)
(10000, 2) (10000, 1)

```

```
In [42]: 1 p9.ggplot(XOR_data,p9.aes(x='X0',y='X1',color='p')) + p9.geom_point()
```



```
Out[42]: <ggplot: (-9223371947463327495)>
```

Exercise 2. (15 pts) Create a neural net with dense hidden layers (including intercepts for each hidden unit) and logistic loss function. Train it on the above dataset with 2 dimensional input using Adam and 32 minibatch size. You may want to increase the number of iterations until you see convergence. You can assess visually by plotting the predictions and comparing to the true probability.

1. Try ReLu activation as well as sigmoid activation for 1 hidden layer and 3 units.
2. Try ReLu activation with 2 hidden layers and 3 units each.
3. Try (2) but with any combination of ReLu and Sigmoid activation.


In [83]:

```

1  #1. Relu
2  sess = tf.reset_default()
3
4  def create_model():
5      # create inputs
6      input_ph = tf.placeholder(dtype=tf.float32, shape=[None, 2])
7      output_ph = tf.placeholder(dtype=tf.float32, shape=[None,1])
8
9      # create variables
10     W0 = tf.get_variable(name='W0', shape=[2,3], initializer=tf.contrib.layers.xavier_in
11     #W1 = tf.get_variable(name='W1', shape=[20, 20], initializer=tf.contrib.layers.xavie
12     W1 = tf.get_variable(name='W1', shape=[3, 1], initializer=tf.contrib.layers.xavier_i
13
14     b0 = tf.get_variable(name='b0', shape=[3], initializer=tf.constant_initializer
15     #b1 = tf.get_variable(name='b1', shape=[20], initializer=tf.constant_initializer(0.)
16     b1 = tf.get_variable(name='b1', shape=[1], initializer=tf.constant_initializer
17
18     weights = [W0,W1]#[W0, W1, W2]
19     biases = [b0,b1]#[b0, b1, b2]
20     activations = [tf.nn.relu,tf.nn.sigmoid]#[tf.nn.sigmoid, tf.nn.relu, tf.nn.sigmoid]
21
22     # create computation graph
23     layer = input_ph
24     print('input shape = ',input_ph.get_shape())
25     for W, b, activation in zip(weights, biases, activations):
26         layer = tf.matmul(layer, W) + b
27         if activation is not None:
28             layer = activation(layer)
29     output_pred = layer
30     print(output_ph, output_pred)
31     return input_ph, output_ph, output_pred
32
33     input_ph, output_ph, output_pred = create_model()
34
35     # create loss
36     #product = tf.multiply(output_ph,output_pred)
37     #logloss = tf.reduce_sum(-tf.log(1 + tf.exp(-(product))))
38     #logloss = -tf.reduce_sum((tf.multiply(output_ph,tf.log(output_pred)) + tf.multiply((1 -
39     logloss = tf.losses.log_loss(output_ph,output_pred)
40
41     # create optimizer
42     opt = tf.train.AdamOptimizer().minimize(logloss)
43
44     # initialize variables
45     sess.run(tf.global_variables_initializer())
46     # create saver to save model variables
47     saver = tf.train.Saver()
48
49     # run training
50     batch_size = 32
51     for training_step in range(50000):
52         #print('in the training')
53         # get a random subset of the training data
54         indices = np.random.randint(low=0, high=len(inputs), size=batch_size)
55         input_batch = inputs[indices]
56         output_batch = outputs[indices]
57
58         # run the optimizer and get the mse
59         _, logloss_run = sess.run([opt, logloss], feed_dict={input_ph: input_batch, output_p
60
61         # print the mse every so often
62         if training_step % 1000 == 0:
63             print('{0:04d} logloss: {1:.3f}'.format(training_step, logloss_run))
64             saver.save(sess, '/tmp/model.ckpt')

```

```
In [82]: 1 sess = tf_reset_default()
2
3 # create the model
4 input_ph, output_ph, output_pred = create_model()
5
6 # restore the saved model
7 saver = tf.train.Saver()
8 saver.restore(sess, "/tmp/model.ckpt")
9
10 output_pred_run = sess.run(output_pred, feed_dict={input_ph: inputs})
11
12 XOR_data_copy = XOR_data.copy()
13 XOR_data_copy.iloc[:,2] = output_pred_run
14 p9.ggplot(XOR_data_copy,p9.aes(x='X0',y='X1',color='p')) + p9.geom_point()
input shape = (?, 2)
Tensor("Placeholder_1:0", shape=(?, 1), dtype=float32) Tensor("Sigmoid:0", shape=(?, 1), dtype=float32)
INFO:tensorflow:Restoring parameters from /tmp/model.ckpt
```



```
Out[82]: <ggplot: (-9223371947452406952)>
```

In [84]:


```

1  #1. Sigmoid
2  sess = tf.reset_default()
3
4  def create_model():
5      # create inputs
6      input_ph = tf.placeholder(dtype=tf.float32, shape=[None, 2])
7      output_ph = tf.placeholder(dtype=tf.float32, shape=[None,1])
8
9      # create variables
10     W0 = tf.get_variable(name='W0', shape=[2,3], initializer=tf.contrib.layers.xavier_in
11     #W1 = tf.get_variable(name='W1', shape=[20, 20], initializer=tf.contrib.layers.xavie
12     W1 = tf.get_variable(name='W1', shape=[3, 1], initializer=tf.contrib.layers.xavier_i
13
14     b0 = tf.get_variable(name='b0', shape=[3], initializer=tf.constant_initializer
15     #b1 = tf.get_variable(name='b1', shape=[20], initializer=tf.constant_initializer(0.)
16     b1 = tf.get_variable(name='b1', shape=[1], initializer=tf.constant_initializer
17
18     weights = [W0,W1]#[W0, W1, W2]
19     biases = [b0,b1]#[b0, b1, b2]
20     activations = [tf.nn.sigmoid,tf.nn.sigmoid]#[tf.nn.sigmoid, tf.nn.relu, tf.nn.sigmoi
21
22     # create computation graph
23     layer = input_ph
24     print('input shape = ',input_ph.get_shape())
25     for W, b, activation in zip(weights, biases, activations):
26         layer = tf.matmul(layer, W) + b
27         if activation is not None:
28             layer = activation(layer)
29     output_pred = layer
30     print(output_ph, output_pred)
31     return input_ph, output_ph, output_pred
32
33     input_ph, output_ph, output_pred = create_model()
34
35     # create loss
36     #product = tf.multiply(output_ph,output_pred)
37     #logloss = tf.reduce_sum(-tf.log(1 + tf.exp(-(product))))
38     #logloss = -tf.reduce_sum((tf.multiply(output_ph,tf.log(output_pred)) + tf.multiply((1 -
39     logloss = tf.losses.log_loss(output_ph,output_pred)
40
41     # create optimizer
42     opt = tf.train.AdamOptimizer().minimize(logloss)
43
44     # initialize variables
45     sess.run(tf.global_variables_initializer())
46     # create saver to save model variables
47     saver = tf.train.Saver()
48
49     # run training
50     batch_size = 32
51     for training_step in range(50000):
52         #print('in the training')
53         # get a random subset of the training data
54         indices = np.random.randint(low=0, high=len(inputs), size=batch_size)
55         input_batch = inputs[indices]
56         output_batch = outputs[indices]
57
58         # run the optimizer and get the mse
59         _, logloss_run = sess.run([opt, logloss], feed_dict={input_ph: input_batch, output_p
60
61         # print the mse every so often
62         if training_step % 1000 == 0:
63             print('{0:04d} logloss: {1:.3f}'.format(training_step, logloss_run))
64             saver.save(sess, '/tmp/model.ckpt')

```



```
In [85]: 1 sess = tf_reset_default()
2
3 # create the model
4 input_ph, output_ph, output_pred = create_model()
5
6 # restore the saved model
7 saver = tf.train.Saver()
8 saver.restore(sess, "/tmp/model.ckpt")
9
10 output_pred_run = sess.run(output_pred, feed_dict={input_ph: inputs})
11
12 XOR_data_copy = XOR_data.copy()
13 XOR_data_copy.iloc[:,2] = output_pred_run
14 p9.ggplot(XOR_data_copy,p9.aes(x='X0',y='X1',color='p')) + p9.geom_point()
input shape = (?, 2)
Tensor("Placeholder_1:0", shape=(?, 1), dtype=float32) Tensor("Sigmoid_1:0", shape=(?, 1), dtype=float32)
INFO:tensorflow:Restoring parameters from /tmp/model.ckpt
```



```
Out[85]: <ggplot: (-9223371947451215100)>
```



In [86]:

```

1  #2. Relu
2  sess = tf.reset_default()
3
4  def create_model():
5      # create inputs
6      input_ph = tf.placeholder(dtype=tf.float32, shape=[None, 2])
7      output_ph = tf.placeholder(dtype=tf.float32, shape=[None,1])
8
9      # create variables
10     W0 = tf.get_variable(name='W0', shape=[2,3], initializer=tf.contrib.layers.xavier_in
11     W1 = tf.get_variable(name='W1', shape=[3, 3], initializer=tf.contrib.layers.xavier_i
12     W2 = tf.get_variable(name='W2', shape=[3, 1], initializer=tf.contrib.layers.xavier_i
13
14     b0 = tf.get_variable(name='b0', shape=[3], initializer=tf.constant_initializer
15     b1 = tf.get_variable(name='b1', shape=[3], initializer=tf.constant_initializer
16     b2 = tf.get_variable(name='b2', shape=[1], initializer=tf.constant_initializer
17
18     weights = [W0,W1,W2]#[W0, W1, W2]
19     biases = [b0,b1,b2]#[b0, b1, b2]
20     activations = [tf.nn.relu,tf.nn.relu,tf.nn.sigmoid]#[tf.nn.sigmoid, tf.nn.relu, tf.n
21
22     # create computation graph
23     layer = input_ph
24     print('input shape = ',input_ph.get_shape())
25     for W, b, activation in zip(weights, biases, activations):
26         layer = tf.matmul(layer, W) + b
27         if activation is not None:
28             layer = activation(layer)
29     output_pred = layer
30     print(output_ph, output_pred)
31     return input_ph, output_ph, output_pred
32
33     input_ph, output_ph, output_pred = create_model()
34
35     # create loss
36     #product = tf.multiply(output_ph,output_pred)
37     #logloss = tf.reduce_sum(-tf.log(1 + tf.exp(-(product))))
38     #logloss = -tf.reduce_sum((tf.multiply(output_ph,tf.log(output_pred)) + tf.multiply((1 -
39     logloss = tf.losses.log_loss(output_ph,output_pred)
40
41     # create optimizer
42     opt = tf.train.AdamOptimizer().minimize(logloss)
43
44     # initialize variables
45     sess.run(tf.global_variables_initializer())
46     # create saver to save model variables
47     saver = tf.train.Saver()
48
49     # run training
50     batch_size = 32
51     for training_step in range(50000):
52         #print('in the training')
53         # get a random subset of the training data
54         indices = np.random.randint(low=0, high=len(inputs), size=batch_size)
55         input_batch = inputs[indices]
56         output_batch = outputs[indices]
57
58         # run the optimizer and get the mse
59         _, logloss_run = sess.run([opt, logloss], feed_dict={input_ph: input_batch, output_p
60
61         # print the mse every so often
62         if training_step % 1000 == 0:
63             print('{0:04d} logloss: {1:.3f}'.format(training_step, logloss_run))
64             saver.save(sess, '/tmp/model.ckpt')

```

```
In [87]: 1 sess = tf_reset_default()
2
3 # create the model
4 input_ph, output_ph, output_pred = create_model()
5
6 # restore the saved model
7 saver = tf.train.Saver()
8 saver.restore(sess, "/tmp/model.ckpt")
9
10 output_pred_run = sess.run(output_pred, feed_dict={input_ph: inputs})
11
12 XOR_data_copy = XOR_data.copy()
13 XOR_data_copy.iloc[:,2] = output_pred_run
14 p9.ggplot(XOR_data_copy,p9.aes(x='X0',y='X1',color='p')) + p9.geom_point()
input shape = (?, 2)
Tensor("Placeholder_1:0", shape=(?, 1), dtype=float32) Tensor("Sigmoid:0", shape=(?, 1), dtype=float32)
INFO:tensorflow:Restoring parameters from /tmp/model.ckpt
```



```
Out[87]: <ggplot: (89402592008)>
```


In [88]:

```

1  #3. Relu,Sigmoid
2  sess = tf.reset_default()
3
4  def create_model():
5      # create inputs
6      input_ph = tf.placeholder(dtype=tf.float32, shape=[None, 2])
7      output_ph = tf.placeholder(dtype=tf.float32, shape=[None,1])
8
9      # create variables
10     W0 = tf.get_variable(name='W0', shape=[2,3], initializer=tf.contrib.layers.xavier_in
11     W1 = tf.get_variable(name='W1', shape=[3, 3], initializer=tf.contrib.layers.xavier_i
12     W2 = tf.get_variable(name='W2', shape=[3, 1], initializer=tf.contrib.layers.xavier_i
13
14     b0 = tf.get_variable(name='b0', shape=[3], initializer=tf.constant_initializer
15     b1 = tf.get_variable(name='b1', shape=[3], initializer=tf.constant_initializer
16     b2 = tf.get_variable(name='b2', shape=[1], initializer=tf.constant_initializer
17
18     weights = [W0,W1,W2]#[W0, W1, W2]
19     biases = [b0,b1,b2]#[b0, b1, b2]
20     activations = [tf.nn.relu,tf.nn.sigmoid,tf.nn.sigmoid]#[tf.nn.sigmoid, tf.nn.relu, t
21
22     # create computation graph
23     layer = input_ph
24     print('input shape = ',input_ph.get_shape())
25     for W, b, activation in zip(weights, biases, activations):
26         layer = tf.matmul(layer, W) + b
27         if activation is not None:
28             layer = activation(layer)
29     output_pred = layer
30     print(output_ph, output_pred)
31     return input_ph, output_ph, output_pred
32
33     input_ph, output_ph, output_pred = create_model()
34
35     # create loss
36     #product = tf.multiply(output_ph,output_pred)
37     #logloss = tf.reduce_sum(-tf.log(1 + tf.exp(-(product))))
38     #logloss = -tf.reduce_sum((tf.multiply(output_ph,tf.log(output_pred)) + tf.multiply((1 -
39     logloss = tf.losses.log_loss(output_ph,output_pred)
40
41     # create optimizer
42     opt = tf.train.AdamOptimizer().minimize(logloss)
43
44     # initialize variables
45     sess.run(tf.global_variables_initializer())
46     # create saver to save model variables
47     saver = tf.train.Saver()
48
49     # run training
50     batch_size = 32
51     for training_step in range(50000):
52         #print('in the training')
53         # get a random subset of the training data
54         indices = np.random.randint(low=0, high=len(inputs), size=batch_size)
55         input_batch = inputs[indices]
56         output_batch = outputs[indices]
57
58         # run the optimizer and get the mse
59         _, logloss_run = sess.run([opt, logloss], feed_dict={input_ph: input_batch, output_p
60
61         # print the mse every so often
62         if training_step % 1000 == 0:
63             print('{0:04d} logloss: {1:.3f}'.format(training_step, logloss_run))
64             saver.save(sess, '/tmp/model.ckpt')

```

```
In [89]: 1 sess = tf_reset_default()
2
3 # create the model
4 input_ph, output_ph, output_pred = create_model()
5
6 # restore the saved model
7 saver = tf.train.Saver()
8 saver.restore(sess, "/tmp/model.ckpt")
9
10 output_pred_run = sess.run(output_pred, feed_dict={input_ph: inputs})
11
12 XOR_data_copy = XOR_data.copy()
13 XOR_data_copy.iloc[:,2] = output_pred_run
14 p9.ggplot(XOR_data_copy,p9.aes(x='X0',y='X1',color='p')) + p9.geom_point()
input shape = (?, 2)
Tensor("Placeholder_1:0", shape=(?, 1), dtype=float32) Tensor("Sigmoid_1:0", shape=(?, 1), dtype=float32)
INFO:tensorflow:Restoring parameters from /tmp/model.ckpt
```



```
Out[89]: <ggplot: (-9223371947446081876)>
```