

```
In [1]: 1 ### I disabled my GPU since most of you will not have that
        2 import os
        3 os.environ["CUDA_DEVICE_ORDER"] = "PCI_BUS_ID" # see issue #152
        4 os.environ["CUDA_VISIBLE_DEVICES"] = ""
```

DavisSML Tensorflow Lab 2

Some content is taken from

- www.tensorflow.org (<http://www.tensorflow.org>)
- <http://rail.eecs.berkeley.edu/deeprcourse/> (<http://rail.eecs.berkeley.edu/deeprcourse/>)

To install tensorflow and keras use pip: `pip install tensorflow`, `pip install keras`

Keras

- high level API, typically do not interact with Tensors as in base tensorflow
- build Models from Layers
- integrated with Datasets API
- extensible : can build custom layers

```
In [1]: 1 from __future__ import print_function
        2 import keras
        3 from keras import layers, datasets, models
        4 from keras import backend as K
        5 import tensorflow as tf
        6 import matplotlib.pyplot as plt
        Using TensorFlow backend.
```

Models and Layers

- Models contain Layers (in a graph)
- Layers contain all of the tensors, variables, operations, etc.
- Sequential graph is just a stack of layers (one feeds into the next)

```
In [2]: 1 model = keras.Sequential()
        2 # Adds a densely-connected layer with 64 units to the model:
        3 model.add(layers.Dense(64, activation='relu'))
        4 # Add another:
        5 model.add(layers.Dense(64, activation='relu'))
        6 # Add a softmax layer with 10 output units:
        7 model.add(layers.Dense(10, activation='softmax'))
```

See the layers that are in the model here:

```
In [9]: 1 model.layers

Out[9]: [<keras.layers.core.Dense at 0x22af7b4e1d0>,
        <keras.layers.core.Dense at 0x22af7b4e3c8>,
        <keras.layers.core.Dense at 0x22af7b4e550>]
```

Layers have

- activation parameter: can be `tf.sigmoid` or 'sigmoid' for example
- kernel initializer, bias initializer: how to initialize coefficients or intercept
- kernel_regularizer, bias_regularizer: regularizers for coefficients and intercepts

```
In [10]: 1 model = keras.Sequential()
2         # Adds a densely-connected layer with 64 units to the model:
3         model.add(layers.Dense(64, activation=tf.sigmoid, kernel_initializer='orthogonal'))
4         # Add another:
5         model.add(layers.Dense(64, activation='relu', kernel_regularizer = keras.regularizers
6         # Add a softmax layer with 10 output units:
7         model.add(layers.Dense(10, activation='softmax'))
```

Model compile:

- optimizer: ex. `tf.train.AdamOptimizer`, `tf.train.RMSPropOptimizer`, `tf.train.GradientDescentOptimizer`
- loss: `mse`, `categorical_crossentropy`, and `binary_crossentropy`
- metrics: monitoring training

Documentation: [Losses \(https://keras.io/losses/\)](https://keras.io/losses/), [Metrics \(https://keras.io/metrics/\)](https://keras.io/metrics/), [Optimizers \(https://www.tensorflow.org/api_docs/python/tf/keras/optimizers\)](https://www.tensorflow.org/api_docs/python/tf/keras/optimizers)

```
In [11]: 1 model.compile(optimizer=tf.train.AdamOptimizer(0.001),
2                   loss='categorical_crossentropy',
3                   metrics=['accuracy'])
```

Model fit will run the optimizer with `epochs`, `batch_size`, `validation_data`

```
In [12]: 1 import numpy as np
          2
          3 data = np.random.random((1000, 32))
          4 labels = np.random.random((1000, 10))
          5
          6 model.fit(data, labels, epochs=10, batch_size=32)
Epoch 1/10
1000/1000 [=====] - 5s 5ms/step - loss: 15.4082 - acc: 0.0880
Epoch 2/10
1000/1000 [=====] - 0s 96us/step - loss: 14.4723 - acc: 0.0750
Epoch 3/10
1000/1000 [=====] - 0s 97us/step - loss: 13.6801 - acc: 0.0950
Epoch 4/10
1000/1000 [=====] - 0s 97us/step - loss: 13.0243 - acc: 0.0800
Epoch 5/10
1000/1000 [=====] - 0s 95us/step - loss: 12.4948 - acc: 0.1130
Epoch 6/10
1000/1000 [=====] - 0s 97us/step - loss: 12.0927 - acc: 0.1070
Epoch 7/10
1000/1000 [=====] - 0s 96us/step - loss: 11.8260 - acc: 0.1040
Epoch 8/10
1000/1000 [=====] - 0s 97us/step - loss: 11.6751 - acc: 0.1150
Epoch 9/10
1000/1000 [=====] - 0s 95us/step - loss: 11.5801 - acc: 0.1100
Epoch 10/10
1000/1000 [=====] - 0s 100us/step - loss: 11.5148 - acc: 0.1110
```

```
Out[12]: <keras.callbacks.History at 0x22adfbe3898>
```

Custom layer

- build: Create the weights of the layer. Add weights with the add_weight method.
- call: Define the forward pass.
- compute_output_shape: Specify how to compute the output shape of the layer given the input shape.

```
In [13]: 1 for l in model.layers:
          2     print("{} : {} to {}".format(l.name, l.input_shape, l.output_shape))
dense_4 : (None, 32) to (None, 64)
dense_5 : (None, 64) to (None, 64)
dense_6 : (None, 64) to (None, 10)
```

```
In [9]: 1 class MyLayer(layers.Layer):
2
3     def __init__(self, output_dim, **kwargs):
4         self.output_dim = output_dim
5         super(MyLayer, self).__init__(**kwargs)
6
7     def build(self, input_shape):
8         shape = tf.TensorShape((input_shape[1], self.output_dim))
9         # Create a trainable weight variable for this layer.
10        self.kernel = self.add_weight(name='kernel',
11                                     shape=shape,
12                                     initializer='uniform',
13                                     trainable=True)
14        # Make sure to call the `build` method at the end
15        super(MyLayer, self).build(input_shape)
16
17    def call(self, inputs):
18        return tf.matmul(inputs, self.kernel)
19
20    def compute_output_shape(self, input_shape):
21        shape = tf.TensorShape(input_shape).as_list()
22        shape[-1] = self.output_dim
23        return tf.TensorShape(shape)
```

Image data

We will create a model that includes

- convolutional layers: convolution of a kernel of certain size, number of channels, activation
- pooling: max pooling for example, pool size
- dropout: dropout probability

```
In [14]: 1 batch_size = 128
2 num_classes = 10
3 epochs = 4
4
5 # input image dimensions
6 img_rows, img_cols = 28, 28
```

```
In [15]: 1 # the data, split between train and test sets
2 (x_train, y_train), (x_test, y_test) = datasets.mnist.load_data()
3
4 if K.image_data_format() == 'channels_first':
5     x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
6     x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
7     input_shape = (1, img_rows, img_cols)
8 else:
9     x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
10    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
11    input_shape = (img_rows, img_cols, 1)
Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz (https://s3.
amazonaws.com/img-datasets/mnist.npz)
11493376/11490434 [=====] - 2s 0us/step
```

```
In [16]: 1 x_train = x_train.astype('float32')
2 x_test = x_test.astype('float32')
3 x_train /= 255
4 x_test /= 255
5 print('x_train shape:', x_train.shape)
6 print(x_train.shape[0], 'train samples')
7 print(x_test.shape[0], 'test samples')
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

```
In [17]: 1 # convert class vectors to binary class matrices
2 y_train = keras.utils.to_categorical(y_train, num_classes)
3 y_test = keras.utils.to_categorical(y_test, num_classes)
4
5 model = keras.Sequential()
6 model.add(layers.Conv2D(32, kernel_size=(3, 3),
7                           activation='relu',
8                           input_shape=input_shape))
9 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
10 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
11 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
12 #model.add(layers.Dropout(0.25))
13 model.add(layers.Flatten())
14 model.add(layers.Dense(128, activation='relu'))
15 model.add(layers.Dropout(0.5))
16 model.add(layers.Dense(num_classes, activation='softmax'))
```


```
In [18]: 1 for l in model.layers:
2     print("{} : {} to {}".format(l.name, l.input_shape, l.output_shape))
conv2d_1 : (None, 28, 28, 1) to (None, 26, 26, 32)
max_pooling2d_1 : (None, 26, 26, 32) to (None, 13, 13, 32)
conv2d_2 : (None, 13, 13, 32) to (None, 11, 11, 64)
max_pooling2d_2 : (None, 11, 11, 64) to (None, 5, 5, 64)
flatten_1 : (None, 5, 5, 64) to (None, 1600)
dense_7 : (None, 1600) to (None, 128)
dropout_1 : (None, 128) to (None, 128)
dense_8 : (None, 128) to (None, 10)
```

```
In [19]: 1 model.compile(loss=keras.losses.categorical_crossentropy,
2               optimizer=keras.optimizers.Adadelta(),
3               metrics=['accuracy'])
4 model.summary()
5
6 model.fit(x_train, y_train,
7         batch_size=batch_size,
8         epochs=epochs,
9         verbose=1,
10        validation_data=(x_test, y_test))
11 score = model.evaluate(x_test, y_test, verbose=0)
12 print('Test loss:', score[0])
13 print('Test accuracy:', score[1])
```

Layer (type)	Output Shape	Param #
=====		
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_1 (Flatten)	(None, 1600)	0
dense_7 (Dense)	(None, 128)	204928
dropout_1 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
=====		
Total params: 225,034		
Trainable params: 225,034		
Non-trainable params: 0		
=====		
Train on 60000 samples, validate on 10000 samples		
Epoch 1/4		
60000/60000 [=====] - 10s 172us/step - loss: 0.2947 - acc: 0.9080 - val_loss: 0.0788 - val_acc: 0.9734		
Epoch 2/4		
60000/60000 [=====] - 7s 124us/step - loss: 0.0886 - acc: 0.9730 - val_loss: 0.0373 - val_acc: 0.9872		
Epoch 3/4		
60000/60000 [=====] - 7s 124us/step - loss: 0.0639 - acc: 0.9805 - val_loss: 0.0336 - val_acc: 0.9881		
Epoch 4/4		
60000/60000 [=====] - 7s 124us/step - loss: 0.0526 - acc: 0.9844 - val_loss: 0.0281 - val_acc: 0.9898		
Test loss: 0.02812100595554948		
Test accuracy: 0.9898		

```
In [20]: 1 plt.imshow(x_train[0,...,0],cmap='Greys')
```

```
Out[20]: <matplotlib.image.AxesImage at 0x22bf69a2748>
```



```
In [21]: 1 layer_outputs = [layer.output for layer in model.layers[:4]]  
2 activation_model = models.Model(inputs=model.input, outputs=layer_outputs) # Creates a model
```

```
In [22]: 1 activations = activation_model.predict(x_train[[0],...])  
2 # Returns a list of five Numpy arrays: one array per layer activation
```


In [23]:

```

1 layer_names = []
2 for layer in model.layers:
3     layer_names.append(layer.name) # Names of the layers, so you can have them as part o
4
5 images_per_row = 16
6
7 for layer_name, layer_activation in zip(layer_names, activations): # Displays the featur
8     n_features = layer_activation.shape[-1] # Number of features in the feature map
9     size = layer_activation.shape[1] #The feature map has shape (1, size, size, n_featur
10    n_cols = n_features // images_per_row # Tiles the activation channels in this matrix
11    display_grid = np.zeros((size * n_cols, images_per_row * size))
12    for col in range(n_cols): # Tiles each filter into a big horizontal grid
13        for row in range(images_per_row):
14            channel_image = layer_activation[0,
15                                           :, :,
16                                           col * images_per_row + row]
17            channel_image -= channel_image.mean() # Post-processes the feature to make i
18            channel_image /= channel_image.std()
19            channel_image *= 64
20            channel_image += 128
21            channel_image = np.clip(channel_image, 0, 255).astype('uint8')
22            display_grid[col * size : (col + 1) * size, # Displays the grid
23                               row * size : (row + 1) * size] = channel_image
24
25    scale = 1. / size
26    plt.figure(figsize=(scale * display_grid.shape[1],
27                        scale * display_grid.shape[0]))
28    plt.title(layer_name)
29    plt.grid(False)
30    plt.imshow(display_grid, aspect='auto', cmap='viridis')

```

c:\users\lahiru d. chamain\anaconda3\envs\tfgpu\lib\site-packages\ipykernel_launcher.py:18: RuntimeWarning: invalid value encountered in true_divide



after this layer it's goof to flatten because there is no neighboring info

Exercise 1 In the above model we selected a sequence of convolutions, pooling, dropouts, and dense layers. Run the following experiments with 4 epochs and 32 minibatch size, each time reporting the accuracy.

- Remove dropout layers.
- Remove first max pooling layer, does it take longer or shorter to train?
- Try sigmoid activation functions instead of ReLU.

```
In [25]: 1 batch_size = 32
          2 num_classes = 10
          3 epochs = 4
```

```
In [29]: 1 # Remove the dropouts and then max pooling -1
2 model = keras.Sequential()
3 model.add(layers.Conv2D(32, kernel_size=(3, 3),
4                         activation='relu',
5                         input_shape=input_shape))
6 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
7 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
8 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
9 #model.add(layers.Dropout(0.25))
10 model.add(layers.Flatten())
11 model.add(layers.Dense(128, activation='relu'))
12 #model.add(layers.Dropout(0.5))
13 model.add(layers.Dense(num_classes, activation='softmax'))
14
15 model.compile(loss=keras.losses.categorical_crossentropy,
16               optimizer=keras.optimizers.Adadelta(),
17               metrics=['accuracy'])
18 model.summary()
19
20 model.fit(x_train, y_train,
21         batch_size=batch_size,
22         epochs=epochs,
23         verbose=1,
24         validation_data=(x_test, y_test))
25 score = model.evaluate(x_test, y_test, verbose=0)
26 print('Test loss:', score[0])
27 print('Test accuracy:', score[1])
```

Layer (type)	Output Shape	Param #
=====		
conv2d_9 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_8 (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_10 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_9 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten_5 (Flatten)	(None, 1600)	0
dense_15 (Dense)	(None, 128)	204928
dense_16 (Dense)	(None, 10)	1290
=====		

Total params: 225,034

Trainable params: 225,034

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/4

60000/60000 [=====] - 17s 282us/step - loss: 0.1232 - acc: 0.9616 - val_loss: 0.0447 - val_acc: 0.9848

Epoch 2/4

60000/60000 [=====] - 16s 267us/step - loss: 0.0396 - acc: 0.9883 - val_loss: 0.0355 - val_acc: 0.9870

Epoch 3/4

60000/60000 [=====] - 16s 268us/step - loss: 0.0275 - acc: 0.9917 - val_loss: 0.0287 - val_acc: 0.9897

Epoch 4/4

60000/60000 [=====] - 16s 271us/step - loss: 0.0207 - acc: 0.9935 - val_loss: 0.0270 - val_acc: 0.9909

Test loss: 0.02702516848493142

Test accuracy: 0.9909

We got a training acc of 99.35% and test acc of 99.09%. Training takes on average $16/\text{epoch} \times 4 \text{ epochs} = 64$ seconds

```
In [27]: 1 # Remove the dropouts
2 model = keras.Sequential()
3 model.add(layers.Conv2D(32, kernel_size=(3, 3),
4                           activation='relu',
5                           input_shape=input_shape))
6 #model.add(layers.MaxPooling2D(pool_size=(2, 2)))
7 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
8 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
9 #model.add(layers.Dropout(0.25))
10 model.add(layers.Flatten())
11 model.add(layers.Dense(128, activation='relu'))
12 #model.add(layers.Dropout(0.5))
13 model.add(layers.Dense(num_classes, activation='softmax'))
14
15 model.compile(loss=keras.losses.categorical_crossentropy,
16               optimizer=keras.optimizers.Adadelta(),
17               metrics=['accuracy'])
18 model.summary()
19
20 model.fit(x_train, y_train,
21           batch_size=batch_size,
22           epochs=epochs,
23           verbose=1,
24           validation_data=(x_test, y_test))
25 score = model.evaluate(x_test, y_test, verbose=0)
26 print('Test loss:', score[0])
27 print('Test accuracy:', score[1])
```

Layer (type)	Output Shape	Param #
=====		
conv2d_7 (Conv2D)	(None, 26, 26, 32)	320
conv2d_8 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_7 (MaxPooling2	(None, 12, 12, 64)	0
flatten_4 (Flatten)	(None, 9216)	0
dense_13 (Dense)	(None, 128)	1179776
dense_14 (Dense)	(None, 10)	1290
=====		

Total params: 1,199,882

Trainable params: 1,199,882

Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/4

60000/60000 [=====] - 36s 596us/step - loss: 0.1133 - acc: 0.9647 - val_loss: 0.0411 - val_acc: 0.9868

Epoch 2/4

60000/60000 [=====] - 35s 582us/step - loss: 0.0344 - acc: 0.9893 - val_loss: 0.0308 - val_acc: 0.9897

Epoch 3/4

60000/60000 [=====] - 35s 582us/step - loss: 0.0213 - acc: 0.9936 - val_loss: 0.0326 - val_acc: 0.9889

Epoch 4/4

60000/60000 [=====] - 35s 584us/step - loss: 0.0132 - acc: 0.9960 - val_loss: 0.0367 - val_acc: 0.9903

Test loss: 0.03671608702375579

Test accuracy: 0.9903

We got a training acc of 99.60% and test acc of 99.03%. Training takes on average $35/\text{epoch} \times 4 \text{ epochs} = 140$ seconds. With the maxpooling-1 removed, we get lower training and test accuracy and takes longer time, x2 times compared to the case with maxpooling-1. This is because with maxpooling it halves the image shape after that layer, making the rest of the propagation faster.

```
In [30]: 1 #Sigmoid instead of relu
2
3 model = keras.Sequential()
4 model.add(layers.Conv2D(32, kernel_size=(3, 3),
5                         activation='sigmoid',
6                         input_shape=input_shape))
7 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
8 model.add(layers.Conv2D(64, (3, 3), activation='sigmoid'))
9 model.add(layers.MaxPooling2D(pool_size=(2, 2)))
10 model.add(layers.Dropout(0.25))
11 model.add(layers.Flatten())
12 model.add(layers.Dense(128, activation='sigmoid'))
13 #model.add(layers.Dropout(0.5))
14 model.add(layers.Dense(num_classes, activation='softmax'))
15
16 model.compile(loss=keras.losses.categorical_crossentropy,
17              optimizer=keras.optimizers.Adadelta(),
18              metrics=['accuracy'])
19 model.summary()
20
21 model.fit(x_train, y_train,
22         batch_size=batch_size,
23         epochs=epochs,
24         verbose=1,
25         validation_data=(x_test, y_test))
26 score = model.evaluate(x_test, y_test, verbose=0)
27 print('Test loss:', score[0])
28 print('Test accuracy:', score[1])
```

Layer (type)	Output Shape	Param #
=====		
conv2d_11 (Conv2D)	(None, 26, 26, 32)	320

max_pooling2d_10 (MaxPooling)	(None, 13, 13, 32)	0

conv2d_12 (Conv2D)	(None, 11, 11, 64)	18496

max_pooling2d_11 (MaxPooling)	(None, 5, 5, 64)	0

dropout_2 (Dropout)	(None, 5, 5, 64)	0

flatten_6 (Flatten)	(None, 1600)	0

dense_17 (Dense)	(None, 128)	204928

dense_18 (Dense)	(None, 10)	1290
=====		
Total params: 225,034		
Trainable params: 225,034		
Non-trainable params: 0		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/4
60000/60000 [=====] - 17s 280us/step - loss: 0.7014 - ac
c: 0.7624 - val_loss: 0.1675 - val_acc: 0.9489
Epoch 2/4
60000/60000 [=====] - 16s 265us/step - loss: 0.1522 - ac
c: 0.9528 - val_loss: 0.0955 - val_acc: 0.9679
Epoch 3/4
60000/60000 [=====] - 16s 268us/step - loss: 0.1066 - ac
c: 0.9669 - val_loss: 0.0776 - val_acc: 0.9745
Epoch 4/4
60000/60000 [=====] - 16s 267us/step - loss: 0.0876 - ac
c: 0.9725 - val_loss: 0.0652 - val_acc: 0.9786
```

We got a training acc of 97.25% and test acc of 97.86%. Training takes on average 16/epoch*4 epochs = 64 seconds